

Debugging Machine Learning Tasks

Aleksandar Chakarov^A, Aditya Nori^B, Sriram Rajamani^B,
Shayak Sen^C, and Deepak Vijaykeerthy^D

^AUniversity of Colorado, Boulder

^BMicrosoft Research

^CCarnegie Mellon University

^DIBM Research

March 24, 2016

Abstract

Unlike traditional programs (such as operating systems or word processors) which have large amounts of code, machine learning tasks use programs with relatively small amounts of code (written in machine learning libraries), but voluminous amounts of data. Just like developers of traditional programs debug errors in their code, developers of machine learning tasks debug and fix errors in their data. However, algorithms and tools for debugging and fixing errors in data are less common, when compared to their counterparts for detecting and fixing errors in code. In this paper, we consider classification tasks where errors in training data lead to misclassifications in test points, and propose an automated method to find the root causes of such misclassifications. Our root cause analysis is based on Pearl’s theory of causation, and uses Pearl’s PS (Probability of Sufficiency) as a scoring metric. Our implementation, PSI, encodes the computation of PS as a probabilistic program, and uses recent work on probabilistic programs and transformations on probabilistic programs (along with gray-box models of machine learning algorithms) to efficiently compute PS. PSI is able to identify root causes of data errors in interesting data sets.

1 Introduction

Machine learning techniques are used to perform data-driven decision-making in a large number of diverse areas including image processing, medical diagnosis, credit decisions, insurance decisions, email spam detection, speech recognition, natural language processing, robotics, information retrieval and online advertising. Over time, these techniques have been honed and tuned, and are now at a stage where machine learning libraries [1, 2] are used as black-boxes by

programmers with little or no expertise in the details of the machine learning algorithms themselves. The black-box nature of the reuse, however, has an unfortunate downside. Current implementations of machine learning techniques provide little insight into why a particular decision was made. Because of this absence of transparency, debugging the outputs of a machine learning algorithm has become incredibly hard.

Most programmers who implement machine learning use libraries to build models from voluminous training data, and then use these models to perform predictions. These machine learning libraries often employ complex, stochastic, or approximate, search and optimization algorithms that search for an optimal model for a given training data set. The model is then applied to a set of unseen test samples in the hope of satisfactory generalization. When generalization fails, i.e., an incorrect result is produced for a test input, it is often difficult to debug the cause of the failure. Such failures can arise due to several reasons. Common causes for failure include bugs in the implementation of the machine learning algorithm, incorrect choice of features, incorrect setting of parameters (such as degree of the polynomial for regression or number of layers in a neural network) when invoking the machine learning library, and noise in the training set. Over time, bugs in implementation of machine learning algorithms get detected and fixed. There is a lot of work in feature selection [3], and parameter choices can be made by systematically building models for various parameter values and choosing the model with the best validation score [4]. However, since training data is typically voluminous, errors in training data are common and notoriously difficult to debug. This suggests a new class of debugging problems where programs (machine learning classifiers) are learnt from data and bugs in a program are now the result of faults in the data.

In this paper, we focus on debugging machine learning tasks in the presence of errors in training data. Specifically we consider classification tasks, which are typically implemented using algorithms such as logistic regression [5] and boosted decision trees [6]. Suppose we train a classifier on training data (which has errors), and the classifier produces incorrect results for one or more test points. We desire to produce an automated procedure to identify the *root cause* of this failure. That is, we would like to identify a subset of training points that influences the classification for these test points the most. Therefore, correcting mistakes in these training points is most likely to fix the incorrect results.

Our algorithm for identifying root causes is inspired by the structural equations framework of causation, as formulated by Judea Pearl [7, 8]. We think of each of the training data points as possible causes of the misclassification in the test data set, and calculate for each such training point, a score corresponding to how likely it is that the current label for that point is the cause for the misclassification of the test data set. A simple measure of the score of a training point can be obtained by merely flipping the label of the training point and observing if the flip improves the results of the classifier on test points. However, such a simple measure does not work when errors exist in several training points, and several training points *together* cause the incorrect results in the test points. Thus, the score we calculate for each training point t considers

alternate *counterfactual worlds*, where training points are labeled with several possible values (other than the value in the training data), and sums up the probability that flipping the label of t causes the misclassification error in the test data, among all such alternate worlds. In Pearl’s framework, such a score is called the *probability of sufficiency* or PS for short.

One of the main difficulties in calculating the probability of sufficiency is that the classifier (or model) needs to be relearned for alternate worlds. Each of these model computing steps (also called as *training* steps) is expensive. We use a “gray box” view of the machine learning library, and profile key intermediate values (that are hand-picked for each machine learning algorithm) during the initial training phase. Using these values, we build a gray-box abstraction of the training process by which the model for a new training set (which is obtained by flipping certain number of training labels) can be obtained efficiently without the need to perform complete (and expensive) retraining. Finally, we are able to amortize the cost of computing the PS score by sharing common work *across* the computation for different training points.

In order to carry out these optimizations, we model the PS computation as a *probabilistic program* [9]. Probabilistic programs allow us to represent all of the above optimizations such as using gray-box models, using instrumented values from actual training runs, and sharing work across multiple PS computations as program transformations. We are also able to leverage recent progress in efficient inference of probabilistic programs to scale the computation of PS scores to large data sets.

We have implemented our root cause detection algorithm in a tool PSI. PSI currently works with two popular classifiers: (1) logistic regression, and (2) boosted decision trees. For these classifiers, PSI runs a production quality implementation of the techniques, profiles specific values and builds an abstract gray-box model of the classifier, which avoids expensive re-training. Armed with this gray-box model, PSI performs scalable inference to compute the PS values for all points in the training set. PSI is able to identify root causes of misclassifications in several interesting data sets. In summary, the main contributions of this paper are as follows:

- We propose using the structural equations framework of causality, and specifically Pearl’s PS score to compute root causes of failures in machine learning algorithms.
- We model the PS computation as a probabilistic program, and this enables us to leverage efficient techniques developed to perform inference on probabilistic programs to calculate PS scores. We build gray-box models of the machine learning techniques by profiling actual training runs of the library, and using profiled values to build abstract models of the training process. We amortize work across PS computations of different training points. Probabilistic programs allow us to carry out these optimizations and reason about them as program transformations.
- We have built a tool PSI implementing the approach for logistic regression

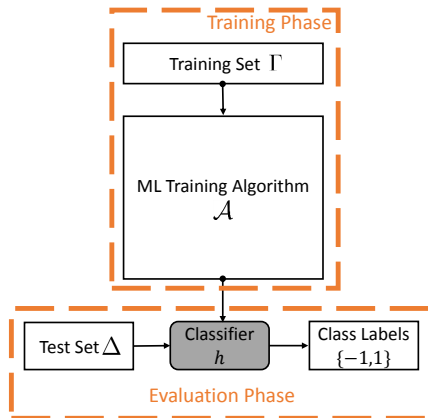


Figure 1: A two-stage design flow of a machine learning task: *training phase* in which the ML algorithm \mathcal{A} is applied to training set Γ to learn classifier h , and *evaluation phase* to judge the quality of h on test set Δ .

and boosted decision trees. PSI is able to identify root causes of misclassifications in several interesting data sets. We hypothesize that this approach can be generalized to other machine learning tasks as well.

2 Overview

We motivate our approach through the experience of Alice, a typical developer who uses machine learning.

2.1 Typical Scenario

Alice is not a machine learning expert, but needs to write a classifier for images of vehicles and animals. Mallory is a machine learning expert who built a classification library using state of the art machine learning techniques. Alice decides to use Mallory’s library, and since machine learning libraries are driven by data, she carefully collects some amount of training data $\{x^i\}_i$ with images of cats, dogs, elephants trucks, cars, buses etc., with labels $y^i = -1$ or $y^i = 1$, stating whether an image is that of a vehicle or an animal respectively. She partitions it into a *training set* $\Gamma = \{(x^i, y^i)\}_{i=1}^M$, and a *test set* Δ , and picks out her favorite ML algorithm, logistic regression, to learn a binary classifier that separates vehicles from animals.

Alice runs Mallory’s impeccable implementation of logistic regression on her training set Γ to learn the classifier $h : \Gamma \rightarrow \{-1, 1\}$. She then evaluates h over the test set Δ . This two-stage design flow is common, and is shown in Figure 1. Alice is happy with most classifications being correct, but unhappy that a particular image x_Δ of a small car has been incorrectly classified as an

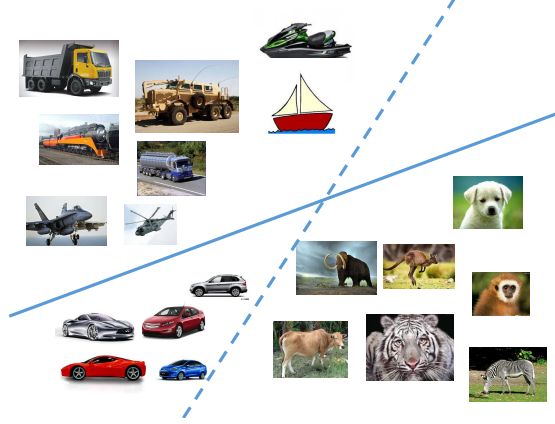


Figure 2: A classification example. All cars are incorrectly labeled as animals in the training set. Thus the learned classifier (shown as a solid line) incorrectly classifies cars in the test data set as animals. If the errors in the training set are fixed, then a correct classifier (shown in dotted line) will be learnt. Our goal is to identify these errors in the training set efficiently.

animal. She wants to find out an explanation for why $h(x_\Delta) \neq y^t$. She suspects some training data may be mislabeled, so she wants to know what set of training instances $\Gamma_{x_\Delta} \subseteq \Gamma$ “caused” x_Δ to be misclassified.

The cause for misclassification in her case is that each of the 5 cars in her training data are classified as animals due to an error in her script which collected the training data, as shown in Figure 2. Since her test instance is a car, the classifier incorrectly classifies it as an animal.

Alice’s tried-and-tested approach to debugging is to start at the point of the error and work backward and find which portions of the code led to the error [10]. Logistic regression uses the training data to find a model θ which is an n -dimensional vector, where n is the number of features in the training data. Once the model θ is calculated (using training data), on the test input x_Δ , the output is given by $h(x_\Delta) = \frac{1}{1+e^{-\theta^T \cdot x_\Delta}}$, which is a sigmoid function applied to the dot product $\theta^T \cdot x_\Delta$ (see Sections 5.1 for more details). Thus, the output for the test point x_Δ depends on *all* the indices of the θ vector, and it turns out that θ depends on *all* the elements of the training set. Thus, unfortunately, tracing backward through program dependences does not enable Alice to narrow down the root cause of the failure.

Next, Alice turns to experimental approaches to answering why the classifier $h(\cdot)$ classifies the car x_Δ as an animal [11]. She picks training data points at random, and changes their labels and looks at the outcome after rerunning logistic regression. She finds that changing the label of a single training point makes no difference to the classifier. Hence, she chooses subsets of training

points, changes all their labels simultaneously, reruns the training phase, and observes if the resulting (changed) classifier correctly classifies the car x_Δ . She finds that switching the label on sets of training data points that include several cars has much greater influence on the classification of x_Δ than other sets of points. Given that combinations of causes are involved, she wonders if there is a systematic way to identify all (or most of) the possible causes. Each black-box experiment Alice runs is time consuming, since training is expensive. Thus, Alice wonders if there is a better way to perform her experiments more systematically and more efficiently.

2.2 Our Approach

Our goal is to identify root causes by systematically and efficiently carrying out the experimentation approach that Alice attempts to do. There are two main difficulties that Alice faces in her experimentation: (1) Changing any single training instance does not fix the error. On the other hand, trying to identify subsets of training points which fix the error is infeasible due to the explosively large number of possible subsets of points. (2) Each experiment takes a long time to run, since it reruns the training algorithm from scratch. We consider each of these difficulties in detail.

In order to counter the first difficulty, instead of considering sets of training points, we measure influence of individual points in causing the classification. We use Pearl’s theory of probabilities of causation [7] which formalizes the influence which a training instance t has on x_Δ being misclassified by considering all possible labellings of the training points as possible worlds, and measuring the number of worlds in which changing the label of t changes the outcome of classification of x . In this spirit, we perform experiments with a large number of alternate worlds, and for each world w we can evaluate if in world w , changing the label of t influences the classification of x_Δ . Given a world w , we perform this experiment simultaneously for all training points, and record which of the training points influence the classification of x_Δ in world w . We repeat the entire experiment with several alternate worlds, and compute an aggregate score for each training point t on x over all these alternate worlds. Such an aggregate score is called Probability of Sufficiency or PS (see Section 3.2 for more details).

The second difficulty is the cost of building a model for each experiment (also known as training). Black-box experimentation does not scale because there are a huge number of training points, and running the entire implementation of the machine learning algorithm for each experimental trial takes a lot of time. We can potentially take a white-box view since we have access to its source; however, incorporating the full details of the source poses a scalability issue to even the state-of-art inference techniques. As a practical compromise, we take a gray-box view of the machine learning implementation, and profile key intermediate values (that are hand-picked for each machine learning algorithm) during the initial training phase. Using the profiling information, we are able to use intermediate values from runs of one training set Γ to efficiently calculate the PS value for a related training set Γ' , where Γ and Γ' differ in a small

number of labels. The exact values profiled, and the nature of the gray-box model depends on the specific machine learning algorithm. In Section 5, we show gray-box models for both logistic regression and decision trees.

Probabilistic Program for PS computation. We model the computation of the PS score as a *probabilistic program* [9]. Writing down the PS computation as a probabilistic program enables us to use techniques developed for probabilistic program inference to calculate PS. It also enables to concretely represent and reason about various optimizations to calculate PS as program transformations. The probabilistic program we write for the PS score (see Section 3.3) models a Bernoulli distribution with mean given by the PS value, and inference over the probabilistic program gives us an estimate of the PS value.

Empirical Results. We have implemented our approach in the tool PSI. In Section 6, we show empirical evaluation of PSI using synthetic benchmarks as well as real-world data sets. We introduce 10% random errors in the training data for synthetic benchmarks, and systematic 10% errors in the training data for real-world data sets. We find that PSI is able to identify a significant fraction of the systematically introduced errors using the calculated PS scores. For randomly introduced errors, the performance of PSI depends on the nature of the benchmark. For instance, if the benchmark already has a lot of inherent noise, then adding extra noise does not really change the classifier, and hence there is not enough information to perform root causing. We also evaluate the effect of changing the labels for training points with high PS scores on the validation score. We find that the validation score improves monotonically as we make changes for points with high PS scores, and starts to degrade if we change labels for PS scores below 0.28, confirming that higher PS scores are more likely candidates for root causing errors. We also find that using PS values from multiple test points enables PSI to find more root causes, as well as improve validation score. See Section 6 for more details.

3 Preliminaries

In this section, we lay out the formal details of the machine learning tasks we consider. We introduce the intuition and formalization of the probabilistic causality framework we propose. We also introduce probabilistic programs, our specification mechanism for probabilistic causality.

3.1 Machine Learning Applications

We consider classification tasks in supervised machine learning [5]. Formally, a *machine learning task* $(\Gamma, \Delta, \mathcal{A})$ consists of a training set $\Gamma = \{(x^i, y^i)\}_{i=1}^M$, and a test set $\Delta = \{(x^j, y^j)\}_{j=1}^N$ of labeled samples (x^j, y^j) , where $x^i, x^j \in \mathbb{R}^n$ are called *feature vectors*, $y^i, y^j \in \{-1, 1\}$ are called classification *labels*, and \mathcal{A} is

an ML classification algorithm. Generally, developing a machine learning task consists of two phases as shown in Figure 1:

1. *Training Phase*, in which the training algorithm \mathcal{A} is applied to the training set Γ to derive as output a binary classification function $h : \mathbb{R}^n \rightarrow \{-1, 1\}$ that maps each feature vector $x^i \in \mathbb{R}^n$ to a classification label -1 or 1 .
2. *Evaluation Phase*, in which h is applied to each sample d in Δ . The result of this phase is an evaluation score S that is defined as follows:

$$S \equiv \sum_{(x^d, y^d) \in D} I(h(x^d) \neq y^d),$$

where $I(\varphi)$ is the 0-1-indicator function.

The goal of a machine learning task is to compute a classification function h over Γ that generalizes well to the unseen test set Δ so as to minimize the evaluation score S .

3.2 Probabilistic Counterfactual Causality

Before we formalize the definition of causality, we motivate it with an example.

Example 1. Consider a scenario where people vote for a government; choices are A or B . In the final count of votes, out of 101 people who voted, 56 voted for A and 45 for B . Everyone who voted for A contributed to A 's winning; however, no individual voter v appears to be the sole cause (affects the outcome of the election). The theory of causality via counterfactuals [8, 7] proposes a solution to this apparent predicament by considering alternate worlds in which an individual v 's vote indeed affects the outcome. Informally, v 's vote is a cause if there exists an alternate world (in this case an assignment to other voters' choices) such that in this alternate world, the outcome is decided by v 's vote alone. Such alternate worlds are called *counterfactuals* (counter to the actual world we observed).

While considering the existence of counterfactual worlds helps establish causality in a qualitative sense, we find it useful to consider a quantitative measure of causation. For example, consider the US presidential election where different states have different numbers of electoral votes. California has 55 electoral votes, whereas Wyoming has only 3 electoral votes. Intuitively, California voting for A is a *more significant* cause for A winning than Wyoming voting for A . This notion can be captured by considering the number of counterfactual worlds in which the outcome is affected. The number of alternate worlds in which California affects the outcome of the election (any world in which the difference in votes, barring California's, is less than 55) is greater than the number of alternate worlds in which Wyoming affects the outcome of the election. We next focus on formally stating this definition.

3.2.1 Causal Models

Informally, a *causal model*¹ is a relationship between the inputs to any system and its output. A simple model for the voting example with n voters is to have input variables Y_1, \dots, Y_n (one for each voter), where each $Y_i \in \{-1, 1\}$, denoting a vote for A or B respectively. Assume each vote Y_i carries weight a_i , the causal model relating the inputs to the output (the predicate “Has A won?”) is therefore:

$$a_1 Y_1 + \dots + a_n Y_n > 0$$

Definition 1. A *causal model* M is a set of random *input variables* $\{Y_1, \dots, Y_n\}$, each taking values in domain D , an output variable X over some output domain D' , and a *structural equation* $X = f(Y_1, \dots, Y_n)$ relating the output variable X to the input variables, through some well-defined (deterministic) function $f : D^n \rightarrow D'$. An assignment w to all input variables is called a *world*.

To consider counterfactual possibilities, we need to express quantities of the following form: “The value X would have obtained, in the world w , had Y been y ”. Our tool for expressing counterfactuals is an intervention. An *intervention* $[Y_i \leftarrow y_i]$ is a substitution of y_i for the value of Y_i in a world w . Formally, we represent $f(w[Y \leftarrow y])$ using the notation $X_{[Y \leftarrow y]}^w$.

To reason about uncertainty about the inputs to the causal model, we augment causal models with a probability measure over worlds.

Definition 2. A *probabilistic causal model* (M, p) is a causal model M augmented with a probability measure p over the input variables, i.e., a function from $D^n \rightarrow [0, 1]$, where D is the domain of each of the n input variables.

As the structural equations are deterministic, p also defines a distribution over the output X of the causal model M , where:

$$Pr(X = x) = \sum_{\{w | f(w) = x\}} p(w)$$

Similarly, we can define the probabilities of counterfactual statements:

$$Pr(X_{[Y \leftarrow y]} = x) = \sum_{\{w | X_{[Y \leftarrow y]}^w = x\}} p(w)$$

We can also define the probabilities of conditional counterfactual statements as

¹Our presentation is a simplification of the structural model semantics of causality that is suitable for our application. We refer the reader to [8] for a comprehensive introduction to counterfactual causality.

follows:

$$\begin{aligned}
& Pr(X_{[Y \leftarrow y]} = x \mid X \neq x, Y \neq y) \\
&= \frac{Pr(X_{[Y \leftarrow y]} = x, X \neq x, Y \neq y)}{Pr(X \neq x, Y \neq y)} \\
&= \sum_{\{w \mid X_{[Y \leftarrow y]}^w = x\}} p(w \mid X \neq x, Y \neq y)
\end{aligned} \tag{1}$$

This is the conditional probability of “ X being x when Y is y , given X is not x and Y is not y ”. This conditional probability may appear to be zero at first glance, but under the counterfactual interpretation, X and $X_{[Y \leftarrow y]}$ are actually evaluated under different assignment for the input variable Y . In fact, this conditional probability, known as *probability of sufficiency*, is the quantitative measure we use to measure causality in our setting.

Probability of Sufficiency. Assume that in the world $w : Y_1 = y_1, \dots, Y_n = y_n$, it is the case that $X = x$. Then the probability of sufficiency (PS) of some $Y_i = y_i$ being the cause for $X = x$ is defined as:

$$PS(Y_i) \equiv Pr(X_{[Y_i \leftarrow y_i]} = x \mid X \neq x, Y_i \neq y_i) \tag{2}$$

According to equation (1), the probability of sufficiency measures the probability of each world in which $X \neq x$ and $Y_i \neq y_i$, but changing Y_i to its true value y_i results in X changing to x . For the voting example, where each Y_i represents an individual vote, and X represents the outcome of the election, these worlds are ones in which an individual vote affects the outcome of the election. Therefore, the probability of sufficiency for a voter n who voted for A , is the probability that $a_1 Y_1 + \dots + a_{n-1} Y_{n-1} + a_n > 0$ given $a_1 Y_1 + \dots + a_{n-1} Y_{n-1} - a_n < 0$. Intuitively, for a larger number of electoral votes a_n , the set of worlds that satisfy the above condition is greater, thereby leading to a greater *PS* value for California with 55 electoral votes than Wyoming with 3 electoral votes.

3.3 Probabilistic Programs

Probabilistic programs [9] are “usual” programs with two additional constructs: (1) a *sample* statement, that provides the ability to draw values from distributions, and (2) an *observe* statement that provides the ability to condition on the values of variables. The purpose of a probabilistic program is to implicitly specify a probability distribution. *Probabilistic inference* is the problem of computing an explicit representation of the probability distribution implicitly specified by a probabilistic program.

Consider the probabilistic program shown in Figure 3. This program tosses two fair coins (simulated by draws from a Bernoulli distribution with mean 0.5 in lines 1 and 2), and assigns the outcomes of these coin tosses to the Boolean variables c_1 and c_2 . The *observe* statement in line 3 blocks all executions of the

```

TwoCoins()
1:   $c_1 := \text{sample}(\text{Bernoulli}(0.5));$ 
2:   $c_2 := \text{sample}(\text{Bernoulli}(0.5));$ 
3:   $\text{observe}(c_1 \vee c_2);$ 
4:   $\text{return}(c_1, c_2);$ 

```

Figure 3: A simple probabilistic program.

```

 $PS_i(\mathcal{A}, \Gamma, \varphi)$ 
1:   $\Gamma.Y := \text{sample}(p_T);$ 
2:   $\text{observe}(\Gamma.Y_i \neq y^i);$ 
3:   $c := \mathcal{A}(\Gamma);$ 
4:   $\text{observe}(\neg\varphi(c));$ 
5:   $\Gamma.Y_i := y^i;$ 
6:   $c' := \mathcal{A}(\Gamma);$ 
7:   $\text{return}(\varphi(c'));$ 

```

Figure 4: A probabilistic program that specifies $PS(Y_i)$.

program that do not satisfy the condition specified by it (the Boolean expression $(c_1 \vee c_2)$). The meaning of a probabilistic program is the probability distribution over the expression returned by it. For our example program *TwoCoins*, the distribution specified is the distribution over the pair (c_1, c_2) : $Pr(c_1 = 0, c_2 = 1) = Pr(c_1 = 1, c_2 = 0) = Pr(c_1 = 1, c_2 = 1) = 1/3$, and $Pr(c_1 = 0, c_2 = 0) = 0$.

4 A Probabilistic Program to Compute PS

In this section, we show how to encode the computation of the PS value as a probabilistic program. Formally, we are given:

1. A machine learning task $(\Gamma, \Delta, \mathcal{A})$, with training set $\Gamma = \{(x^i, y^i)\}_{i=1}^N$, test set $\Delta = \{(x^j, y^j)\}_{j=1}^M$ and machine learning algorithm \mathcal{A} .
2. An error $\varphi(c)$, which is a predicate on the classifier c produced by \mathcal{A} (when run on training data Δ).

An example of a predicate $\varphi(c)$ is $c(x^t) \neq y^t$, where $(x^t, y^t) \in \Delta$. Another example predicate is that the total number of misclassifications over the entire test set Δ is less than some threshold value. For each training instance $(x^i, y^i) \in \Gamma$, we wish to measure the probability of sufficiency of the label y^i for causing the error $\varphi(c)$.

We first define a causal model for $\varphi(c)$. With each training label y^i , we associate an input random variable Y_i of the model. Thus, lower-case variables will y^i denote instances (or samples), and upper case variables Y_i the corresponding

random variable. The probability distribution $p_T(Y_1, \dots, Y_N)$ represents prior beliefs over the training labels, and is based on the training set Γ . The structural equation f is simply the machine learning algorithm \mathcal{A} .

A succinct way of representing $PS(Y_i)$ is given by the probabilistic program shown in Figure 4. The notation $\Gamma.Y$ represents the labels in the training set Γ . Line 1 reassigns the labels in Γ to a new sample of labels from the distribution p_T . Each assignment to $\Gamma.Y$ represents a different world (see Definition 1). The observe statements in lines 2 and 4 reflect the counterfactual essence of the probability of sufficiency. They correspond to the conditions $Y \neq y$ and $X \neq x$ in Equation (1). Line 5 models the intervention where Y_i is set to y_i , and line 6 reruns the learning algorithm \mathcal{A} with the changed training set (where Y_i has been updated). The return value of this probabilistic program is a Bernoulli random variable that expresses precisely the quantity $PS(Y_i)$ as defined in Equation (2). Specifying $PS(Y_i)$ as a probabilistic program allows us to leverage recent advances in inference techniques for probabilistic programs, as well as apply various program transformations to enable efficient and scalable inference.

5 Implementing Computation of PS

In Section 4, we showed that computation of PS can be encoded as a probabilistic program. We can thus implement computing PS using recent advances in inference techniques for probabilistic programs [12, 13]. These techniques estimate the posterior probability distribution of probabilistic programs using sampling. However, directly performing inference for the probabilistic program in Figure 4 is intractable for the following reasons:

1. We potentially need to consider all subsets (2^N) of a large number N of training points. Even for modest ML tasks, N is typically greater than 1000.
2. Moreover, Figure 4 requires inference to be performed once for each label Y_i , where the number of training labels can be very large.
3. For each subset, inference needs to be performed over the full implementation details of the training algorithm \mathcal{A} . Handling such highly optimized C++ code (usually several hundred lines) is beyond the scalability of any existing sampling technique for probabilistic inference. Finally, \mathcal{A} needs to be re-run for each relabeling.

We tackle these three challenges by employing the following assumptions and techniques:

- *Low noise level* - We assume that only a small number of training points are mislabeled. This is encoded in the distribution p_T (in line 1 in Figure 4) which decays fast as the number of mislabeled points grows. We can thus restrict our samples to a small portion of the space of all possible

```

 $PS_i(\mathcal{A}, \Gamma, \varphi)$ 
1:  $\Gamma.Y := \text{sample}(p_T);$ 
2:  $c := \mathcal{A}(\Gamma);$ 
3:  $\text{observe}(\neg\varphi(c));$ 
4: if  $(\Gamma.Y_i \neq y^i)$ 
5:    $\Gamma.Y_i := y^i;$ 
6:    $c' := \mathcal{A}(\Gamma);$ 
7:    $\text{return}(\varphi(c'));$ 
8:  $\text{return}(0);$ 

```

Figure 5: An efficient probabilistic program that specifies $PS(Y_i)$.

mislabeled subsets. This is implemented in the probabilistic program by choosing a prior distribution p_T that is heavily biased toward the given training set Γ .

- *Sample reuse* - To address challenge 2, we observe that samples drawn for one label y_i can be reused for other labels. We can think about this as the following transformation on probabilistic programs: (1) move the observe statement from line 2 in Figure 4 after the observe statement in line 4; and (2) rewrite that observe statement as a conditional (shown in Figure 5). Even though the programs in Figure 4 and Figure 5 are equivalent, the program in Figure 5 performs an important optimization. Since the lines 1–3 in Figure 5 are the same for all i , the work for executing these statements can be shared across all i . Then, for each sample generated, after a quick check that condition in line 4 is satisfied, we execute (in parallel for all such i) lines 5–7.
- *Model approximation & Robustness* - Instead of operating over the complex implementation of training algorithm \mathcal{A} , we approximate it suitably via an approximate causal model $\hat{\mathcal{A}}$. The key observation is that some ML algorithms are generally robust [14]. When a small subset of the input changes, the internal computations leading to the final classifier do not change significantly. In fact, being insensitive to small changes in inputs is desirable for machine learning algorithms to reduce overfitting. Thus, we design approximate causal models for two widely used classification algorithms logistic regression (see Section 5.1) and decision trees (see Section 5.2), which enable efficient computation of $PS(Y_i)$ below.

Next, we present the essential details of our approximate causal models for Logistic Regression (Section 5.1) and Boosted Decision Trees (Section 5.2).

In Figure 5, we need to run the entire algorithm \mathcal{A} each time constraint $\varphi(c)$ is to be evaluated for a sample. We now describe our approximate models $\hat{\mathcal{A}}$ which are simpler equations for computing $\varphi(c)$ in terms of the training labels $\{Y_i\}_i$. The details on deriving the constraints and general details on these algorithms can be found in the supplementary materials.

5.1 Approximate Logistic Regression Model

Logistic Regression (LR) is a popular statistical classification model that uses a sigmoid scoring function:

$$h(Z) = \frac{1}{1 + e^{-Z}}$$

where $Z = \theta^T x$ is the product of inferred classifier weights θ and the feature vector x of some instance. When logistic regression is used for binary classification, instances with score $h(\theta^T x) < 0.5$ are classified $c(x) = -1$ and those with score $h(\theta^T x) \geq 0.5$ as $c(x) = 1$.

The classifier θ is learnt through an iterative gradient descent process that iteratively finds a better classifier θ to fit the data. For Stochastic Gradient Descent, the classifier is improved by iterating over the following update equation:

$$\theta^K = \theta^{K-1} + \alpha^{K-1} \sum_{i=1}^N y^i x^i h(y^i (\theta^{K-1})^T x^i), \quad (3)$$

Here, α is known as the step size, and N is the number of training instances. The above vector equation represents an update for each component of θ^K .

By making a robustness assumption that score on training points in the penultimate iteration $K - 1$ of gradient descent, i.e. $h((\theta^{K-1})^T x^i)$ does not change greatly, it turns out that the following equation provides an approximation for the final classifier θ on a different labelling of training labels $\{Y^i\}_i$ (details are found in A.1 of supplementary materials).

$$\theta = \theta^{K-1} + \alpha^{K-1} \sum_{i=1}^N Y^i x^i h(y^i (\theta^{K-1})^T x^i)$$

Using this, we can simplify the condition $\varphi(c)$ of the classifier c incorrectly classifying a test point x as follows:

$$\varphi(c) : h_\theta(x) \geq \frac{1}{2} \iff \theta^T x \geq 0, \text{ and} \quad (4)$$

$$h_\theta(x) < \frac{1}{2} \iff \theta^T x < 0, \quad (5)$$

where $\theta^T x$ is simply the product:

$$(\theta^{K-1} + \alpha^{K-1} \sum_{i=1}^N Y^i x^i h(y^i (\theta^{K-1})^T x^i))^T x$$

Notice now that due to this simplification, the condition $\varphi(c)$ is a linear constraint on the input random variables $\{Y_i\}_i$. To compute the coefficients of each Y_i , we need to profile θ^{K-1} and the step size α^{K-1} .

5.2 Approximate Decision Tree Model

Decision trees are tree-shaped statistical classification models in which each internal tree node represents a decision split on a feature value, and associated with each leaf node, is score s . The score $s(x)$ of a particular point x is computed by evaluating the decision at each internal node and taking the corresponding branches until a leaf is reached. The leaf score is returned as the score $s(x)$. Each leaf can be viewed as a region R in the feature space and a tree is a partitioning $\{R_1, \dots, R_L\}$. In Gradient Boosted Decision Trees [15], instead of a single tree, a number of trees called an *ensemble* are iteratively learnt through a gradient descent process. Informally, after n trees have been learnt, the $(n+1)^{\text{th}}$ tree is learnt by fitting a new decision tree on the *error residual* \bar{y}_i of each training label y^i . The error residual can be thought of as the difference in the aggregate score due to n decision trees learnt so far, from the true label. This iterative process stops after a fixed number of learning rounds.

Under a robustness assumption for scores for training points and the learnt regions for each of the decision trees, it turns out that for an alternate assignment $\{Y_i\}_i$ to the training labels, the score $s(x)$ for a particular test point x can be written as follows (see supplementary material for details):

$$s(x) = s_0(x) + \sum_{n=1}^L \sum_{k=1}^{K_n} \sum_{i=1}^N \frac{(Y_i - y_i) \sigma I(x_i \in R_{nk}) I(x \in R_{nk})}{\sum_{x_i \in R_{nk}} |\bar{y}_{ni}| (2\sigma - |\bar{y}_{ni}|)}$$

Here, L is the number of decision trees learnt in the ensemble, K_n is the number of leaves in the n^{th} tree and N is the number of training instances. R_{nk} is the region represented by the k^{th} leaf of the n^{th} tree and \bar{y}_{ni} is the error residual for the i^{th} training label after n iterations.

While this might appear to be a very complicated computation, notice that it is only a linear computation on $\{Y_i\}_i$ and all the coefficients can be precomputed from a single run of the algorithm. As a result, the condition $\varphi(c)$ of test point x being misclassified can be written as $s(x) \geq 0$ (resp. < 0), which is a linear constraint on the training labels Y_i .

Another important implication of this equation is that the coefficient of label Y_i is a multiple of $I(x_i \in R_{nk}) I(x \in R_{nk})$. Therefore, the coefficient of Y_i , is nonzero only when the corresponding training point x_i and the test point x belong to the same leaf in some tree n . Hence, only a small number of training labels Y_i actually have any affect on the classification of a training point x .

5.3 The PSI debugging tool.

We have implemented the computation of PS value using the above mentioned optimizations in the tool PSI. PSI takes a machine learning task $(\Gamma, \Delta, \mathcal{A})$, a set of misclassified test points (bugs) $B \subseteq \Delta$ as input, and produces root causes in the form of training labels as output.

Currently PSI supports logistic regression and decision trees, and uses an approximate model $\hat{\mathcal{A}}$ of these machine learning algorithms as described above.

PSI runs a single execution of the industrial strength implementation of these machine learning techniques. We have hand instrumented the implementation of both algorithms to profile for specific parameters that make \mathcal{A} and $\hat{\mathcal{A}}$ operate in lockstep.

For logistic regression, we record the iterates of θ and line search steps α in the stochastic gradient descent variant of search algorithm. Additionally, we memoize classification scores $h(\theta^T x)$ and reuse across different training labels Y^i .

For boosted decision trees, we record the regions corresponding each individual tree in the ensemble and error residuals \bar{y}_i for each iteration of gradient descent.

The core engine driving PSI is the probabilistic program together with its sampling based inference engine (as defined in Figure 5). This program uses the label distribution $p(\Gamma)$ over the training set Γ , the approximate model $\hat{\mathcal{A}}$ for the training algorithm \mathcal{A} in order to compute the PS scores for the training instances.

Training instances are sorted and highest ranked PS score instances $T_{\varphi(c)} \subseteq T$ are returned as proposed causes. The user then examines the instances in $T_{\varphi(c)}$ closely, fixes any labeling inconsistency they detect, and re-runs \mathcal{A} on the modified set T . In Section 6, we show that flipping the labels for training instances as-is without inspection provides improvements in classifier performance.

t

6 Evaluation

In this section, we empirically evaluate the effectiveness of PSI. All experiments were performed on a system with a 1.8 GHz Intel Xeon processor and 64 GB RAM running Microsoft Windows 8.1.

We evaluate the applicability of PSI with respect to two different debugging requirements: (1) identifying errors in training data, and (2) reducing errors on unseen data. To evaluate these two metrics, we follow the workflow described in Figure 6, where we first add noise to 10% of the training labels of a dataset. This perturbation introduces new misclassifications in the test set. We run PSI on the new misclassifications with the goal of finding the most likely training instances to cause the new misclassifications, and make the following measurements:

1. **Identifying Relevant Noise.** We measure the fraction of training instances that PSI returns that are known to have erroneous labels.
2. **Reducing Validation Error.** We introduce a separate *validation set* of instances independent of the training and test sets. We measure the accuracy of (i) a classifier learnt on the noisy training set, versus (ii) a classifier learnt on the noisy training set with the most probable causes suggested by PSI flipped. We then measure the reduction in errors in the second classifier with respect to the first.

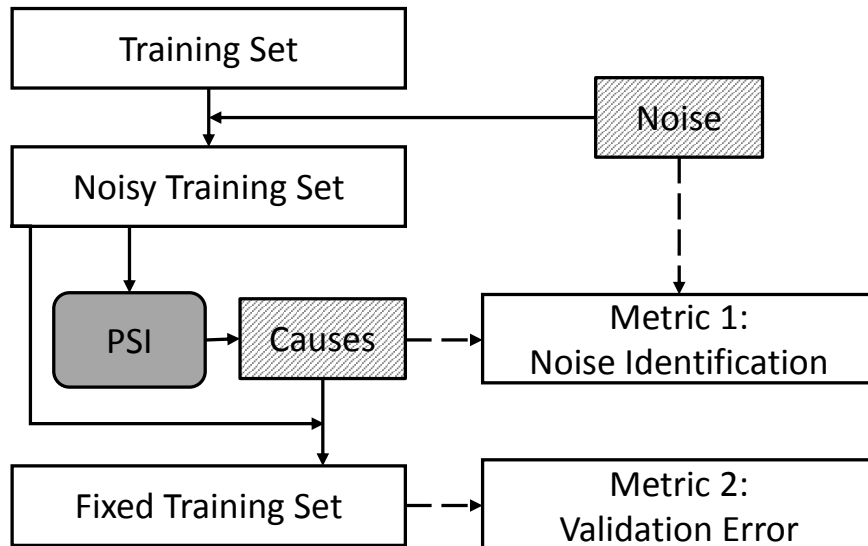


Figure 6: Experimental setup for evaluation.

Table 1 shows the datasets we study. We consider two kinds of datasets: (1) real-world data, and (2) synthetic data. The first two rows of the table are real-world datasets: **sentiment**, an IMDB movie review dataset [16] used for sentiment analysis, and **income**, a census income dataset [17] used for predicting income levels. The third and fourth rows are synthetic datasets: **2gauss** and **concentric**, both produced by generating data from spherical Gaussian distributions.

We also consider two kinds of noise or errors in the datasets: (1) systematic noise, and (2) random noise. We find that PSI can identify systematic noise with significant accuracy, thereby leading to a reduction in classification errors on unseen data. We also observe that adding random noise does not lead to a significant increase in errors, and therefore our causal analysis does not have any room to identify errors. The results of these experiments are summarized in Table 2, and explained in detail in Sections 6.1 (Systematic Noise), and 6.2 (Random Noise). Additionally, we observe that combining information from multiple misclassifications leads to better noise identification.

6.1 Systematic Noise

We add systematic noise to the **sentiment** and **income** datasets, simulating systematic errors in the data collection process. For **sentiment**, we pick a word that appears in 10% percent of instances, and mark these as negative reviews. For **income**, we pick a demographic that covers 10% of the population, and mark them as high income.

Dataset	Features	Training	Noise	Source
sentiment	13132	2000	Syst. – 10%	Real
income	122	4000	Syst. – 10%	Real
2gauss	2	1000	Rand. – 10%	Synth.
concentric	2	2000	Rand. – 10%	Synth.

Table 1: Summary of the benchmark datasets showing the number of features, training instances, noise profile added, and source of each dataset.

Dataset	Accuracy		Validation Error		Time
	LR	DT	LR	DT	
sentiment	0.58	1.00	0.14 → 0.24 → 0.19	0.18 → 0.26 → 0.21	24m
income	0.70	1.00	0.19 → 0.25 → 0.23	0.18 → 0.24 → 0.19	58m
2gauss	0.93	0.38	0.01 → 0.03 → 0.02	0.01 → 0.02 → 0.02	12m
concentric	0.11	0.18	0.47 → 0.49 → 0.48	0.09 → 0.10 → 0.10	14m

Table 2: Summary of experimental results. For each dataset, we report the accuracy (which is the fraction of errors introduced that are identified by PSI), , change in validation error for both logistic regression (LR) and decision trees (DT) and corresponding PSI runtimes in minutes. Validation errors are for a sequence of classifiers learnt on noise free dataset, noisy dataset, noisy dataset with PSI’s suggestions flipped.

When noise is added, the validation error increases. For instance, in the **sentiment** dataset, with logistic regression, the validation error increases from 0.14 to 0.24. PSI is able to successfully identify systematic noise in datasets. For instance, in the **sentiment** dataset, with logistic regression, PSI returns as causes, 58% are points that were incorrectly labeled in the dataset. We call this number accuracy. For decision trees the accuracy is 100%! The accuracy results for **income** benchmark are even better—70% with logistic regression and 100% with decision trees. Note that for datasets with 10% noise, the baseline for randomly picking noise would yield an accuracy of 10%. We discuss the reasons behind the difference in accuracy for the two algorithms in Section 6.3.1.

For systematic errors on a real world dataset, we find that using PSI’s output to flip training labels leads to a reduction in validation error even when not all the points suggested are true errors. For instance, for **sentiment** and the **income** benchmarks, with decision trees, validation error reduces by 0.05 and 0.05 respectively (which is a reduction of 20%).

6.2 Random Noise

In the two synthetic datasets **2gauss** and **concentric**, we randomly flip the labels for 10% of the points as noise to the dataset. We use synthetic datasets to evaluate random noise, and rule out any preexisting systematic noise. We find that for random noise, for a very cleanly separable dataset such as **2gauss**

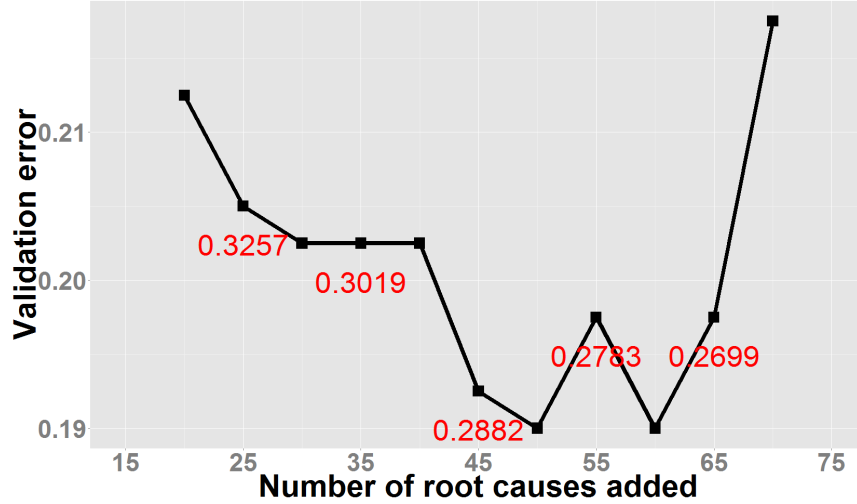


Figure 7: Graph showing reduction in validation error for Logistic Regression as the number of root causes added increases, together with PS threshold τ (in red) for that level.

(note that the initial validation error for this benchmark is 0.01 with both classification schemes, a very low number), it is possible to identify noise with reasonable accuracy (93% for logistic regression, 34% for decision trees).

However, with an inherent noise such as in the **concentric** dataset (note that the initial validation error logistic regression is 0.47), added noise cannot be distinguished from existing noise using causal measures. We observe that validation error changes very little when random noise is added to the dataset, and very few new misclassifications are introduced. This explains why our causal analysis, which depends on observing changes in outcome when inputs are perturbed, does not identify random noise correctly. However, in the presence of very low inherent noise, random noise “stands out” and is easy to identify as is the case in **2gauss**.

6.3 Insights

As shown in Figure 8, we observe combining information from multiple misclassified test points leads to better noise identification. We combine information by choosing $\varphi(c) = \varphi_1(c) \wedge \dots \wedge \varphi_k(c)$, to be the predicate in our PS computation, where each $\varphi_k(c)$ is an individual misclassification. For **sentiment**, precision improves from 0.45 to 0.54 by adding 18 test points. We find this encouraging—with more evidence of test errors, PSI is able to do better root causing, which indicates the robustness of the analysis.

Next, we study how to interpret the PS score, and specifically what (absolute or relative) value of the PS score indicates an error in the training point. We




Figure 8: Graph showing the improvement in noise identification for Logistic Regression runs on `sentiment` dataset as more test points are added.

study this by sorting training points based on the PS score, and measuring the validation error as we flip more points starting with the point with the highest PS score and iteratively flipping more points. Figure 7 shows an interesting trend: validation error reaches a minimum at some PS threshold τ , and changing labels below τ has a detrimental effect. Plotting such a variation between validation error and PS threshold as part of the debugging process could provide an empirical method for choosing optimal PS thresholds.

6.3.1 Logistic Regression vs. Decision Trees

To explain why logistic regression and decision trees behave so differently under our causal analysis, it is important to understand how training points influence the score for a particular test point. Logistic regression fits a single hyperplane that best separates the training set. Therefore, changing the label of any training point affects the hyperplane, and hence the score, but only to a small degree. Decision trees, on the other hand, divide the feature space into smaller regions, one for each tree in the ensemble. To compute a score for a test point $\delta \in \Delta$, only training points that lie in the same region as δ contribute to the score of the test point. In effect, the actual set of points that affect the outcome for a test point is much smaller for decision trees than for logistic regression.

7 Related Work

We briefly review related work in the two most relevant areas: software debugging and counterfactual theory of causality.

7.1 Software Debugging

Debugging software failures is a very well-studied problem. In this section, we survey and compare existing software debugging techniques with PSI.

Program analysis. Program slicing [10] computes a subset of statements in a program that can influence a variable at any program point. Since slicing can be used to effectively reduce the size of the code under analysis, it has many applications in software debugging [18]. Unfortunately, as described in Section 2.1, slicing is not very effective for debugging machine learning tasks as the “data slice” that influences any test error is usually the whole training set. In contrast, PSI is based on a formal notion of causality introduced by [8, 7], and uses this effectively to isolate training data slices that are responsible for test errors.

[19] use a software model checker to generate correct and incorrect traces, and compute the differences between these traces to localize software defects. Unlike their work which looks for bugs in code, we look for bugs in training data. However, they consider both correct and incorrect traces, while we consider only misclassified test data. There may be a way to also make use of correctly classified test data, and we leave this idea for future work.

Delta-Debugging. This debugging technique analyzes differences between failing and passing runs of a program to detect reasons for the failure [20, 21]. In this setting, a cause is defined to be the smallest part of a program state or the smallest part of the input, which when changed, converts a passing run into a failing run. As discussed in Section 2, the main challenge in our setting is that there are usually several causes (several mislabels in the training data) for a misclassified test point, and it is not clear how to use delta-debugging to search through all possible subsets of misclassified training labels. Instead, we use Pearl’s PS score, which is very unique to our work.

Statistical debugging. The CBI project and its variants [22, 23, 24] use information collected from program runs together with a statistical analysis in order to compute predicates that are highly correlated with failures. This technique also requires information from a large number of passing and failing runs of the program, and is applicable when there are bugs in code, and a large number of users are using the same code, and triggering the same bug. In contrast, we assume that the machine learning code is implemented correctly, and the bugs are in the training data. In addition, each user has their own training data, which could have errors, which is a very different problem setting than the one addressed by statistical debugging.

7.2 Counterfactual Causality

PSI computes the *probability of sufficiency* (PS) of a training point in causing classification errors. PS, suggested by Pearl in [7], is part of a rich body of coun-

terfactual theories of causation first proposed by Lewis [25]. In particular, in his book [8], Pearl formulates structural equations as a mathematical framework for reasoning about causality.

Actual Causation. An application of the structural equations framework is Halpern and Pearl’s definition of *actual causation* [26] (known as the HP-Definition). As in our setting, actual causation aims to identify which input variables *actually* caused some outcome. While the HP-definition is able to explain a large variety of subtle issues around causality, operationally, it is intractable to employ in our scenario, since to verify whether an alternate world satisfies the definition, one needs consider an exponentially large number of possibilities. Additionally, the HP-definition is qualitative, and does not suggest a quantitative measure of causality, which is necessary for PSI to rank most likely causes.

8 Conclusion and Future Work

Unlike debuggers for coding errors that are ubiquitous, debuggers for data errors are less common. We consider machine learning tasks, and in particular, classification tasks where incorrect classifiers can be inferred due to errors in training data. We have proposed an approach based in Pearl’s theory of causation, and specifically Pearl’s PS (Probability of Sufficiency) score to rank training points that are most likely causes for having arrived at an incorrect classifier. While the PS score is easy to define, it is expensive to compute. Our tool PSI employs several optimizations to scalably compute PS scores including modeling the computation of the PS score as a probabilistic program, exploiting program transformations and efficient inference techniques for probabilistic programs, and building gray-box models of machine learning algorithms (to save the cost of retraining). Due to these optimizations, PSI is able to correctly root cause data errors in interesting data sets.

Our work opens up several opportunities for interesting future work. An immediate next step is to consider ML tasks involving algorithms such as support vector machines and neural networks [5]. In order to do this, we need to design and develop scalable approximate models for these algorithms.

Another interesting direction is to support a wider class of causes. Potential causes include identifying which feature in a test set causes a misclassification, overfitting in the learning algorithm, incorrect parameters to the learning algorithm, and others.

A practical dimension to explore is scale. We have been able to study datasets with thousands of points using PSI. Industrial big-data systems have millions of points, and more work needs to be done to scale root cause analysis techniques to work at such scale.

References

- [1] M. Hall, E. Frank, G. Holmes, B. Pfahringer, P. Reutemann, and I. H. Witten, “The WEKA data mining software: an update,” *SIGKDD Explorations*, vol. 11, no. 1, pp. 10–18, 2009.
- [2] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [3] I. Guyon and A. Elisseeff, “An introduction to variable and feature selection,” *Journal of Machine Learning Research*, vol. 3, pp. 1157–1182, 2003.
- [4] J. Bergstra and Y. Bengio, “Random search for hyper-parameter optimization,” *Journal of Machine Learning Research*, vol. 13, pp. 281–305, 2012.
- [5] C. M. Bishop *et al.*, *Pattern recognition and machine learning*, vol. 1. springer New York, 2006.
- [6] C. J. Burges, “From ranknet to lambdarank to lambdamart: An overview,” Tech. Rep. MSR-TR-2010-82, June 2010.
- [7] J. Pearl, “Probabilities of causation: Three counterfactual interpretations and their identification,” *Synthese*, vol. 121, no. 1-2, pp. 93–149, 1999.
- [8] J. Pearl, *Causality: Models, Reasoning and Inference*. New York, NY, USA: Cambridge University Press, 2nd ed., 2009.
- [9] A. D. Gordon, T. A. Henzinger, A. V. Nori, and S. K. Rajamani, “Probabilistic programming,” in *Future of Software Engineering (FOSE)*, pp. 167–181, 2014.
- [10] M. Weiser, “Program slicing,” in *International Conference on Software Engineering (ICSE)*, pp. 439–449, 1981.
- [11] A. Zeller, *Why Programs Fail: A Guide to Systematic Debugging*. Morgan Kaufmann, Oct. 2005.
- [12] M. D. Hoffman and A. Gelman, “The no-U-turn sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo,” *Journal of Machine Learning Research*, vol. in press, 2013.
- [13] A. V. Nori, C.-K. Hur, S. K. Rajamani, and S. Samuel, “R2: An efficient mcmc sampler for probabilistic programs,” in *AAAI Conference on Artificial Intelligence (AAAI)*, AAAI, July 2014.
- [14] O. B. A. Elisseeff, “Algorithmic stability and generalization performance,” in *Advances in Neural Information Processing Systems 13: Proceedings of the 2000 Conference*, vol. 13, p. 196, MIT Press, 2001.

- [15] J. H. Friedman, “Greedy function approximation: A gradient boosting machine,” *Annals of Statistics*, vol. 29, pp. 1189–1232, 2000.
- [16] B. Pang, L. Lee, and S. Vaithyanathan, “Thumbs up?: Sentiment classification using machine learning techniques,” in *Proceedings of the ACL-02 Conference on Empirical Methods in Natural Language Processing - Volume 10*, EMNLP ’02, (Stroudsburg, PA, USA), pp. 79–86, Association for Computational Linguistics, 2002.
- [17] D. Newman, S. Hettich, C. Blake, and C. Merz, “Uci repository of machine learning databases,” 1998.
- [18] M. Weiser, “Programmers use slices when debugging,” *Commun. ACM*, vol. 25, no. 7, pp. 446–452, 1982.
- [19] T. Ball, M. Naik, and S. K. Rajamani, “From symptom to cause: localizing errors in counterexample traces,” in *Principles of Programming Languages (POPL)*, pp. 97–1–5, 2003.
- [20] A. Zeller, “Isolating cause-effect chains from computer programs,” in *Proceedings of the 10th ACM SIGSOFT symposium on Foundations of software engineering*, pp. 1–10, ACM, 2002.
- [21] A. Zeller, “Yesterday, my program worked. today, it does not. why?,” in *Software EngineeringESEC/FSE99*, pp. 253–267, Springer, 1999.
- [22] B. Liblit, A. Aiken, A. X. Zheng, and M. I. Jordan, “Bug isolation via remote program sampling,” in *Proceedings of the ACM SIGPLAN 2003 Conference on Programming Language Design and Implementation 2003, San Diego, California, USA, June 9-11, 2003*, pp. 141–154, 2003.
- [23] B. Liblit, M. Naik, A. X. Zheng, A. Aiken, and M. I. Jordan, “Scalable statistical bug isolation,” in *Proceedings of the ACM SIGPLAN 2005 Conference on Programming Language Design and Implementation, Chicago, IL, USA, June 12-15, 2005*, pp. 15–26, 2005.
- [24] T. M. Chilimbi, B. Liblit, K. K. Mehra, A. V. Nori, and K. Vaswani, “HOLMES: effective statistical debugging via efficient path profiling,” in *31st International Conference on Software Engineering, ICSE 2009, May 16-24, 2009, Vancouver, Canada, Proceedings*, pp. 34–44, 2009.
- [25] D. Lewis, “Causation,” *Journal of Philosophy*, vol. 70, no. 17, pp. 556–567, 1973.
- [26] J. Y. Halpern and J. Pearl, “Causes and explanations: A structural-model approach: Part 1: Causes,,” in *UAI* (J. S. Breese and D. Koller, eds.), pp. 194–202, Morgan Kaufmann, 2001.

A Appendix

In this section we present the details of two popular Machine Learning algorithms, Logistic Regression (Section A.1) and Gradient Boosted Decision Trees (Section A.2) as a reference to the reader. The details are standard and can be found in any Machine Learning textbook [5]. We focus on deriving the constraints used in the approximate models of the algorithms found in the main portion of the paper (Sections 5.1 and 5.2).

A.1 Logistic Regression (LR)

Logistic Regression (LR) is a popular statistical classification model based on the sigmoid scoring function:

$$h(Z) = \frac{1}{1 + e^{-Z}},$$

where $Z = \theta^T \cdot x$ with an m -dimensional feature vector x for a sample instance γ , and an m -dimensional vector θ of inferred classification parameters. In this work, we consider LR applied to the binary classification problem in which instances with score $h(\theta^T \cdot x) < 0.5$ are classified $c(x) = -1$ and those with score $h(\theta^T \cdot x) \geq 0.5$ as $c(x) = 1$.

Learning a classifier is posed as an optimization problem over the space of θ values. Specifically, the LR's algorithm attempts to find values of θ that maximize the log-likelihood function:

$$L(\theta) = \sum_{(x^l, y^l) \in T} \log h(y^l z^l),$$

where $z^l = \theta^T \cdot x^l$.

The problem then explores a convex landscape of θ vectors, and an iterative procedure for finding θ candidates (at some step K) can be given using the following recursive equation:

$$\theta^K = \theta^{K-1} + \alpha \cdot \nabla L^{K-1},$$

where α is a step-size in m -dimensional space and the i -th component of the gradient ∇L^{K-1} is:

$$\nabla L_i^{K-1} = \sum_{(x^l, y^l) \in T} y^l x_i^l h(-y^l z^l)$$

with $z^l = (\theta^{K-1})^T \cdot x^l$ for iterates.

Viewing the above equation as a structural equation, we see every training point contributes to the gradient computation. This is a problem as it implies a causal dependence from every input training point to the gradient computed at every iteration. Additionally, in practice a complex optimization algorithm

such as Stochastic Gradient Descent (or a variant) is used to explore the search space.

Unfortunately, as we already stated in Section 2.1 computing the classifier weights θ involves recursive constraints in both the features and labels of *all training instances* (details are found in Section A.1). These problems make causal inference infeasible even for state-of-the-art statistical inference techniques. To reduce the complexity of constraints relating training samples to misclassifications $\varphi(c)$ our approximate causal model $\hat{\mathcal{A}}$ for Logistic Regression uses profiling information from the reference library implementation \mathcal{A} of logistic regression.

PSI performs a single run of \mathcal{A} on Γ to records the values of the parameters essential to computing the weights θ inside the sigmoid classifier h : iterates of θ and line search steps α in the stochastic gradient descent variant of the search algorithm. These parameters (denoted by hats) are then used to simplify key constraints relating training points to the internal computations resulting in misclassifications such as:

$$\theta^K = \hat{\theta}^{K-1} + \hat{\alpha} \nabla \hat{L}, \quad (6)$$

where $\hat{\theta}^{K-1}$ is the $(K-1)$ -th iterate computed by the library implementation, and $\hat{\alpha}_i$ is the step size the implementation took in the along the i -component of the gradient:

$$\nabla \hat{L}_i = \sum_{(x^l, y^l) \in T} Y^l x_i^l h(y^l \hat{z}^l).$$

Here $\hat{z}^l = (\hat{\theta}^{K-1})^T \cdot x^l$ includes the profiling constant. Additionally, we memoize classification scores $h(\theta^T x)$ and reuse across computations for different training labels Y^i using the following well-known property of the sigmoid function $h(-Z) = 1 - h(Z)$.

Because the simplification in (6) allows to compute any θ iterate (say, the last one, or iterates upto a fixed depth d) by treating previous iterates as constants, we can write misclassifications constraints in a simple linear form:

$$\varphi(c) : h_\theta(x) \geq \frac{1}{2} \iff \theta^T x \geq 0, \text{ and} \quad (7)$$

$$h_\theta(x) < \frac{1}{2} \iff \theta^T x < 0, \quad (8)$$

where θ is the final value that the library implementation would compute given the previous iterate $\hat{\theta}^{K-1}$. The simplification translates $\varphi(c)$ directly into a linear constraint over the labels of training instances:

$$\begin{aligned}
(7) &\Leftrightarrow \sum_{i=1}^m \left(\hat{\theta}_i^{K-1} x_i - \hat{\alpha}_i x_i \sum_{(X^l, Y^l) \in T} Y^l x_i^l h(y^l \hat{z}^l) \right) \geq 0, \\
(8) &\Leftrightarrow \sum_{i=1}^m \left(\hat{\theta}_i^{K-1} x_i - \hat{\alpha}_i x_i \sum_{(X^l, Y^l) \in T} Y^l x_i^l h(y^l \hat{z}^l) \right) < 0,
\end{aligned}$$

where feature vectors $X^l = x^l$, step size $\hat{\alpha}$ and value of $h(y^l \hat{z}^l)$ are known constants.

A.2 Gradient Boosted Decision Trees

Decision trees are tree-shaped statistical classification models in which each internal node represents a decision split on a feature value; and, associated with each leaf node is a score s . The score s_p of a particular point p in the feature space \mathbb{R}^n is computed by evaluating the decision at each internal node and taking the corresponding branches until a leaf is reached. The leaf score is then returned as score s_p .

Each leaf can be viewed as a region R of \mathbb{R}^n and a tree as a partitioning $\{R_1, \dots, R_L\}$ of the feature space \mathbb{R}^n . Formally, the tree represents a piecewise constant function $h(x) = \sum_k s_k I(x \in R_k)$. Where, $I(\cdot)$ is the 0-1 indicator function and s_j is the score corresponding to region R_j .

Decision trees are learnt by recursively partitioning the training data on a feature value that maximizes some statistic such as *information gain* until either a fixed number of points are left to be classified or a maximum tree height is reached.

Single decision trees are known to be prone to overfitting since with a large enough tree, the model learnt can be very specific to the training data and often each point can be classified correctly. Overfitting is countered by learning a collection of small trees, known as an *ensemble*, and combining the scores associated with each tree in the ensemble.

One popular technique for learning an ensemble of decision trees is *gradient boosting* [15]. In gradient boosting, a simple tree is initially learnt on the training data T . The residual error E_0 of classifications is used as the new target and another tree is build over the training set fit to the residual errors. This gives an iterative process of refining the training errors that is usually bounded by a fixed number of iterations I_{\max} .

As with the logistic regression classifier learning decision trees is framed as a loss minimization problem. A loss function estimates the error on a training point for a given classifier. For a loss function L , gradient descent proceeds as follows: the initial tree $h_0(x)$ is set to a single partition with some constant score, usually 0; after $n - 1$ trees have been learnt, resulting in a classifier $F_{n-1}(x)$, the next tree $h_n(x)$ is fit to the gradient of the loss function at each training point: $-\frac{\partial L}{\partial F_{n-1}}(x_i)$. A popular choice for a loss function for binary classification

is:

$$L(y, F_n(x)) = \log(1 + e^{-2y\sigma F_n(x)}).$$

Here, σ is a fixed constant known as learning rate. The gradients with respect to this loss function is:

$$\bar{y}_{ni} = -\frac{\partial L(y_i, F_{n-1}(x_i))}{\partial F_{n-1}} = \frac{2y_i\sigma}{1 + e^{-2y\sigma F_{n-1}(x)}}$$

A tree $h_n(x)$ is learnt with \bar{y}_{ni} as the targets. The full details of the derivation can be found in [6]; however, the score of the k th region (leaf) in h_n is as follows:

$$s_{nk} = \frac{\sum_{x_i \in R_{nk}} \bar{y}_{ni}}{\sum_{x_i \in R_{nk}} |\bar{y}_{ni}|(2\sigma - |\bar{y}_{ni}|)}$$

The tree learnt is $h_n(x) = s_{nk}I(x \in R_{nk})$ and the tree is added to the existing classifier to obtain a new classifier $F_n(x) = F_{n-1}(x) + h_n(x)$.

The overall score is the sum of scores from each tree. Let L be the total number of trees and K_n be the number of leaves in the n^{th} tree, then the total score for a test point can be written as:

$$s(x) = \sum_{n=1}^L \sum_{k=1}^{K_n} I(x \in R_{nk}) s_{nk}$$

If $s \geq 0$ the point is classified as 1 else -1 .

We now describe our model for approximating the effect of small perturbations in the input training labels for the library implementation of Gradient Boosted Decision Trees. Here we employ the robustness assumption to simplify the setting. We assume that the trees do not change with small perturbations in the input training labels. Now estimate the change in the score of a leaf s_{nk} due to a flip of label y_i to $-y_i$. For the purposes of estimating the scores for the tree $h_n(x)$ we assume that $F_{n-1}(x)$ does not change due to the flip. Therefore, the new value of the gradient is:

$$\bar{y}'_{ni} = \frac{-2y_i\sigma}{1 + e^{2y\sigma F_{n-1}(x)}} = -2y_i\sigma(1 - \frac{\bar{y}_{ni}}{2y_i\sigma})$$

Also we note that $|\bar{y}_{ni}|(2\sigma - |\bar{y}_{ni}|)$ remains unchanged under the same assumptions. Therefore, the new score s'_{nk} due to flipping y_i to $-y_i$ is:

$$\begin{aligned} s'_{nk} &= s_{nk} + \frac{I(x_i \in R_{nk})(\bar{y}'_{ni} - \bar{y}_{ni})}{\sum_{x_i \in R_{nk}} |\bar{y}_{ni}|(2\sigma - |\bar{y}_{ni}|)} \\ &= s_{nk} + \frac{-2y_i\sigma I(x_i \in R_{nk})}{\sum_{x_i \in R_{nk}} |\bar{y}_{ni}|(2\sigma - |\bar{y}_{ni}|)} \end{aligned}$$

We calculate the score updates for each of the leaves in each tree to calculate the total change in score. Therefore the overall score s' for an assignment \vec{Y} to the training labels is:

$$s'(x) = s(x) + \sum_{n=1}^L \sum_{k=1}^{K_n} \sum_{i=1}^N \frac{(Y_i - y_i)\sigma I(x_i \in R_{nk})I(x \in R_{nk})}{\sum_{x_i \in R_{nk}} |\bar{y}_{ni}|(2\sigma - |\bar{y}_{ni}|)}$$

The condition $\varphi(c)$ which is equivalent to $s \geq 0$ or $s < 0$ is therefore a linear constraint on Y_i 's. This is the result we take away from this section and use in Section [5.2](#).